Data Science
Static data analysis

Stéphane Marchand-Maillet

Master en Sciences Informatiques - Semestre d'Automne

Representation spaces
Quantitative space

Corresponds to a matrix providing the association

ELEMENTS

$$X = \quad \begin{array}{c|cccc} & p_1 & p_2 & \cdots & p_M \\ \hline x^1 & \cdots & \cdots & \cdots & \cdots \\ x^2 & \vdots & \ddots & & \\ \vdots & \vdots & & \ddots & \\ x^N & \vdots & & & \ddots \end{array}$$

VARIABLES

- $\star$ $N \times M$ matrix of variables x_i^k
- $\star$ Each column contains vector $\mathbf{x}_i$ describing p_i via N variables

Representation space
Categorical values

How to represent **symbolic** values?

ELEMENTS

$$X = \begin{array}{c|cccc} & p_1 & p_2 & \cdots & p_M \\ \hline x^1 & 0 & 1 & \cdots & 1 \\ x^2 & 1 & \ddots & & \\ \vdots & \vdots & & 0 & \\ x^N & 1 & & & \ddots \end{array}$$

VARIABLES

- ⋆ Measures the **occurrence** of symbole s in p_i
- ⋆ Binary matrix
- ⋆ Eg: *Vector Space Model*

Representation space
Covariance matrix

Relationships between variables

Matrix $C = XX^\mathsf{T}$ is called **contingency table** (for categorical data). It is related to the **covariance matrix** (for centered numerical data)

VARIABLES

$$C = \begin{array}{c|ccc} & x^1 & x^2 & \cdots \quad x^N \\ \hline x^1 & & & \\ x^2 & \ddots & \vdots & \\ \vdots & \cdots & \sum_k x_k^i x_k^j & \cdots \\ x^N & & \vdots & \ddots \end{array}$$

(with VARIABLES labeling the rows)

$\star$ symbolic : $c_{ij} =$ number of elements with **both** symbols i and j

$\star$ numérique : c_{ij} measures the correlation between variables i and j

Representation space
Similarity table / distance matrix

Measures the distance between elements
Matrix D is called the **distance matrix**

		ELEMENTS		
	p_1	p_2	$\cdots$	p_M
p_1				
P_2	$\ddots$	$\vdots$		
$\vdots$	$\cdots$	$d(\mathbf{x}_i, \mathbf{x}_j)$	$\cdots$	
p_M		$\vdots$		$\ddots$

$D = $ (ELEMENTS)

- $\star$ $M \times M$ matrix, can be very large

- $\star$ if d is a metric, D is p.s.d

Inertia

★ Dataset $\Omega = \{\mathbf{x}_i\}$, $i = 1 \ldots, M$, $\mathbf{x} \in \mathbb{R}^N$

★ Centre of mass :

$$\mathbf{g} = \frac{1}{M} \sum_{\mathbf{x} \in \Omega} \mathbf{x}$$

★ Inertia of Ω w.r.t point $\mathbf{a}$:

$$I_{\mathbf{a}} = \sum_{\mathbf{x} \in \Omega} d(\mathbf{x}, \mathbf{a})^2$$

★ Let $I = I_{\mathbf{g}}$ be the inertia of Ω w.r.t its center of mass $\mathbf{g}$

Inertia (2)

⋆ Inertia of Ω w.r.t subspace $\mathcal{F}$:

$$I_{\mathcal{F}} = \sum_{\mathbf{x} \in \Omega} d^2(\mathbf{x}, \mathcal{F})$$

⋆ $\mathrm{Proj}_{\mathcal{F}}(\mathbf{x})$ is the orthogonal projection of $\mathbf{x}$ onto $\mathcal{F}$, then :

$$I_{\mathcal{F}} = \sum_{\mathbf{x} \in \Omega} d^2(\mathbf{x}, \mathrm{Proj}_{\mathcal{F}}(\mathbf{x}))$$
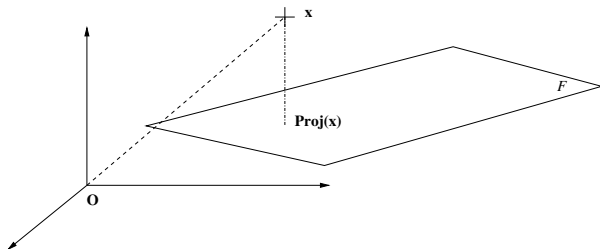
Decomposing inertia (Huygens theorem)

⋆ W.r.t a point:
$$\forall\ \mathbf{a} \in \mathbb{R}^N,\ I_\mathbf{a} = I_\mathbf{g} + d^2(\mathbf{a}, \mathbf{g})$$

→ $\mathbf{g}$ point of **minimum inertia**

⋆ W.r.t a subspace $\mathcal{F}$: given $\mathcal{F}_\mathbf{g}$ vector subspace parrallel to $\mathcal{F}$ via $\mathbf{g}$, then

$$I_\mathcal{F} = I_{\mathcal{F}_\mathbf{g}} + d^2(\mathcal{F}, \mathcal{F}_\mathbf{g})$$

→ $\mathcal{F}_\mathbf{g}$ subspace of **minimum inertia** // to $\mathcal{F}$

Decomposing inertia : Explained inertia

* $\star$ Centered population $\Rightarrow \mathbf{g} = O$, ($O$ origin)
* $\star$ Given $\mathcal{F}$ a vector subspace via O, then

$$I = I_{\mathcal{F}} + I_{\mathcal{F}^{\perp}}, \text{ avec } I_{\mathcal{F}^{\perp}} \text{ inertia of projected points in } \mathcal{F}^{\perp}$$

* $\star$ $I_{\mathcal{F}}$ is called **explained inertia** (by $\mathcal{F}$)
* $\star$ $I_{\mathcal{F}^{\perp}}$ is called **residual inertia** (of $\mathcal{F}$)
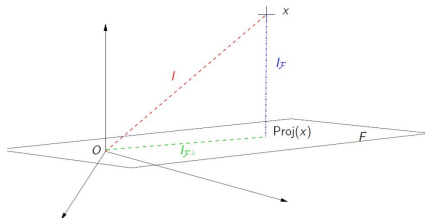
Matrix form

⋆ Centered data $\Rightarrow$

$$I = \sum_{\mathbf{x} \in \Omega} <\mathbf{x}, \mathbf{x}> \Rightarrow I = \sum_{i}^{M} \sum_{j}^{N} (x_i^j)^2 \Rightarrow I = \sum_{j}^{N} \sum_{i}^{M} (x_i^j)^2$$

⋆ if X is the matrix *variables/elements* of size $N \times M$

⋆ then

$$I = \text{trace}(XX^{\mathsf{T}})$$

⋆ matrix XX^{T} is the **inertia matrix**, or $N\times$ the **covariance matrix**

Component analysis

* Multivariate data analysis
  * Ω includes M points defined by N variables x^i
  * $\mathbf{x} \in \mathbb{R}^N$
* We wish to understand the spatial (resp. statistical) distribution of Ω for:
  * Data visualisation
  * Extracting the most important information
* Compression
  * Reconstruction quality
  * Easier handling of the data
* Partition against the data most important properties (features)

Principal Component Analysis (PCA)

The PCA aims at defining, for a population Ω a vector subspace within which the data is represented compactly by uncorrelated variables. These variable will be the new dimension for representing Ω.
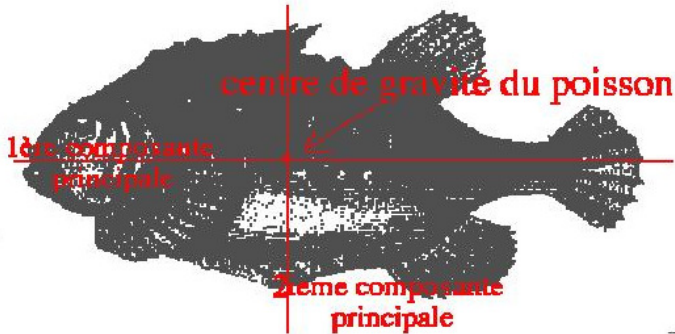
Recall on intertia

- ⋆ Given $\mathbb{R}^N = \Delta_1 \oplus \Delta_2 \oplus \cdots \oplus \Delta_N$ the decomposition of $\mathbb{R}^N$ into 1-D orthogonal subspaces (axes Δ)

- ⋆ Given Ω, the total inertia is decomposed as

$$I = I_{\Delta_1} + I_{\Delta_2} + \cdots + I_{\Delta_N}$$

- ⋆ The PCA searches all subspaces Δ_i s.t $I_{\Delta_i} \geq I_{\Delta_{i+1}}$

- ⋆ The projections onto axes *explaining* the maximum global intertia preserve the maximum of information from the data

- ⋆ if Ω is embedded into a d-dim subspace $\iff I_{\Delta_i} = 0, \ \forall i > d$

Geometric interpretation
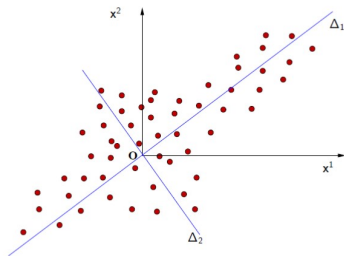
Searching axes with maximum explained inertia

$$\Longleftrightarrow$$

Searching axes supporting maximal variance

Covariance matrix analysis

We search $\mathbf{u}$ minimising the quadratic error (regression)

⋆ $\sum_i ||\mathbf{x}_i - <\mathbf{x}_i, \mathbf{u}> \mathbf{u}||^2 \simeq \sum_i \mathbf{u}^\mathsf{T}\mathbf{x}_i\mathbf{x}_i^\mathsf{T}\mathbf{u} = \textit{Trace}(\mathbf{u}^\mathsf{T}\Sigma\mathbf{u})$

⋆ We add constraint $\mathbf{u}^\mathsf{T}\mathbf{u} = 1$ (to avoid collapse $||\mathbf{u}|| = 0$)

Minimisation

⋆ Lagrange 1D, minimization of
$J = \mathbf{u}^\mathsf{T}\Sigma\mathbf{u} - \lambda(1 - \mathbf{u}\mathbf{u}) \Leftrightarrow \frac{\partial J}{\partial u} = 0 \Leftrightarrow \Sigma\mathbf{u} - \lambda\mathbf{u} = 0$

⋆ $\Leftrightarrow \Sigma\mathbf{u} = \lambda\mathbf{u} \Leftrightarrow \mathbf{u}$ is an eigenvector of Σ

Covariance matrix factorisation

Spectral decomposition

- ⋆ All dimensions are search for simultaneously
- ⋆ ⇔ factorisation of Σ via eigenvalue decomposition

$$\Sigma = U\Lambda U^{\mathsf{T}}, \quad U, \Lambda \in \mathbb{R}^{N \times N}$$

- ⋆ The columns of (rotation matrix) U are the **eigenvectors** $\mathbf{u}_i$ of Σ.
- ⋆ $\Lambda = \mathrm{diag}[\lambda_1, \ldots, \lambda_N]$ (scaling) **eigenvalues** of Σ.

Principal components

* The sorted eigenvalues $\lambda_1 \geq \lambda_2 \geq \cdots \geq \lambda_N$ are the intertia values $I_{\Delta_1} \geq I_{\Delta_2} \geq \cdots \geq I_{\Delta_N}$
* The N sorted eigenvectors $\mathbf{u}_i$ define axes Δ_i and are called the **principal components**
* The new basis for $\mathbb{R}^N$ is now $\{\mathbf{u}_i\}_{i=1,\ldots,N}$
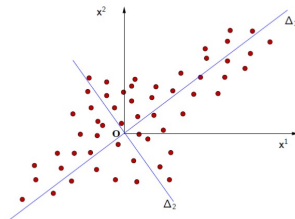
Contribution to total inertia

$$I_{\Delta_i} = \lambda_i, \text{ et } I = \sum_i^N I_{\Delta_i}$$

Definition

1. Absolute contribution of Δ_i to I : $c_{\text{abs}}(\Delta_i/I) = \lambda_i$
2. Relative contribution: $c_{\text{rel}}(\Delta_i/I) = \frac{\lambda_i}{\lambda_1 + \cdots + \lambda_N}$

   $\Rightarrow$ Percentage of inertia explained by Δ_i
3. Percentage for the d first axes :

$$c_{\text{rel}}(\Delta_1 \oplus \Delta_2 \cdots \oplus \Delta_d) = \frac{\lambda_1 + \lambda_2 \cdots + \lambda_d}{\lambda_1 + \lambda_2 \cdots + \lambda_N}$$

Decreasing contribution

Digit data (MNIST)

7291 images 16×16 (8 bits)
$\Rightarrow \mathbf{x}_i \in \mathbb{R}^{256}, \ i = 1 \ldots 7291$

Explained Inertia

CP n°1	CP n°2	CP n°3	CP n°4	CP n°5
$cr(\Delta_1)=18\%$	$cr(\Delta_2)=27\%$	$cr(\Delta_3)=33\%$	$cr(\Delta_4)=39\%$	$cr(\Delta_5)=44\%$

CP n°6	CP n°7	CP n°8	CP n°9	CP n°10
$cr(\Delta_6)=48\%$	$cr(\Delta_7)=51\%$	$cr(\Delta_8)=54\%$	$cr(\Delta_9)=57\%$	$cr(\Delta_{10})=59\%$
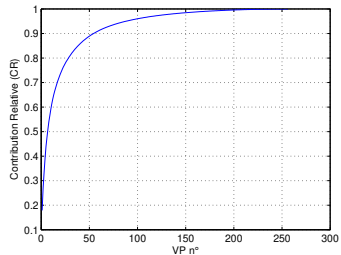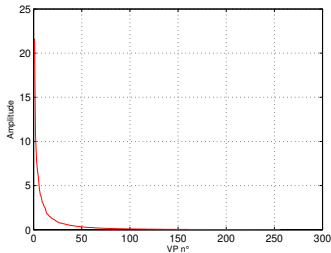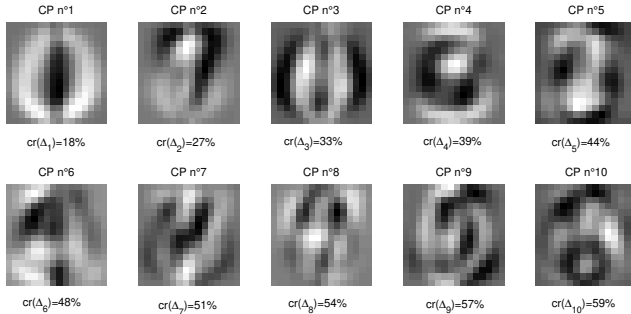
Projection onto principal components

The new space (basis) implies new coordinates for the data

$$y_i^j = <\mathbf{u}_j, \mathbf{x}_i>$$

$\Rightarrow$ the $j$the component of new coordinate $\mathbf{y}_i$ of a data point i is obtained by projecting $\mathbf{x}_i$ onto the jth principal component $\mathbf{u}_j$

$$\mathbf{y}_i = U^\mathsf{T}\mathbf{x}_i$$

Data approximation

If we retain the $d \leq N$ first components (eg $c_{\text{rel}}(\oplus_{i=1\dots d}\Delta_i/I) \geq 90\%$). In this case, the data is approximated in the new space of reduced dimension (subspace $\oplus_{i=1\dots d}\Delta_i$)

$$\tilde{\mathbf{y}}_i^2 = U_d^{\mathsf{T}}\mathbf{x}_i, \quad \tilde{\mathbf{y}} \in \mathbb{R}^d$$

$\Rightarrow U_d \in \mathbb{R}^{N \times d}$ matrix for the dth first components

- $\star$ If $d = 2$ or $3 \rightarrow$ visualisation
- $\star$ If $d \ll N \rightarrow$ compression
- $\star$ Expressivness of the d first components

Data visualisation

$$\tilde{\mathbf{y}}_i^2 = [\mathbf{u}_1^\top \mathbf{x}_i, \mathbf{u}_2^\top \mathbf{x}_i]$$
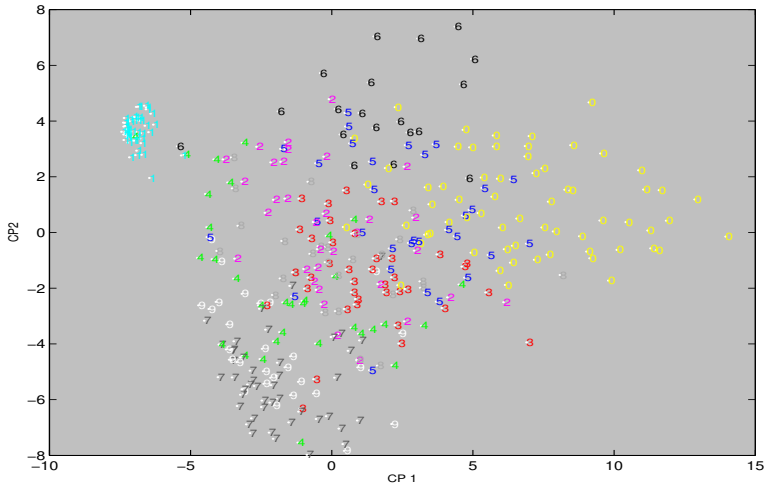
Reconstruction

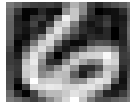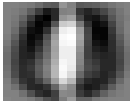$$\tilde{\mathbf{x}}_i = \sum_{j=1}^{d} y_i^j \mathbf{u}_j = {}^t\tilde{\mathbf{y}}_i^d \, U_d$$
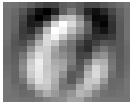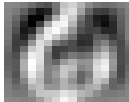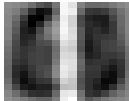

| 4 CP | 16 CP | 64 CP | 256 CP |

| 4 CP | 16 CP | 64 CP | 256 CP |

Representation quality

- ⋆ Given 2 points projected onto Δ_k
  - ○ Far from each other on $\Delta_k \Rightarrow$, far in original space
  - ○ Close from each other on Δ_k, no conclusion. . .
- ⋆ We measure the representation of $\mathbf{x}_i$ on Δ_k by

$$Q_{\Delta_k}(\mathbf{x}_i) = \cos^2(\mathbf{x}_i, \mathbf{u}_k) = \frac{<\mathbf{x}_i, \mathbf{u}_k>^2}{\parallel \mathbf{x}_i \parallel^2}$$

- ⋆ On subspace $E = \Delta_k \oplus \Delta_q \oplus \cdots \oplus \Delta_p$ by

$$Q_E(\mathbf{x}_i) = \cos^2(\mathbf{x}_i, \mathbf{u}_k) + \cos^2(\mathbf{x}_i, \mathbf{u}_q) + \cdots + \cos^2(\mathbf{x}_i, \mathbf{u}_p)$$

Example

Projection on main hyperplane ($\Delta_1 \oplus \Delta_2$)

Data label	0	1	2	3	4	5	6	7	8	9
Quality	0.7	**1.5**	0.4	0.2	0.7	0.2	0.5	0.9	0.4	0.8

Contribution of an element to defining axes

* Absolute contribution of point i to Δ_k

$$c_{\text{abs}}(\mathbf{x}_i, \Delta_k) = \frac{1}{N} < \mathbf{x}_i, \mathbf{u}_k >^2$$
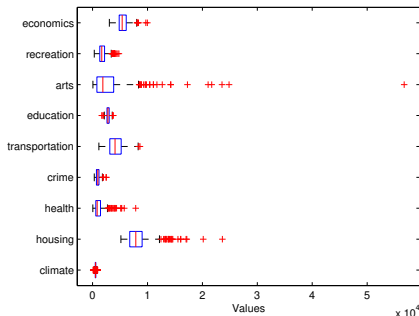
$\Longrightarrow$ The more the projection, the more the point contributes to making the axis exist

* Contribution of a point to the inertia on Δ_k

$$c_{\text{rel}}(\mathbf{x}_i, \Delta_k) = \frac{\text{ca}(\mathbf{x}_i, \Delta_k)}{I_{\Delta_k}} = \frac{< \mathbf{x}_i, \mathbf{u}_k >^2}{\lambda_k}$$

PCA on scaled data

★ Heterogeneous initial variables → validity of linear combination?

★ Example : Evaluation of American cities

★ Scales: depend on categories 100 →> 10000

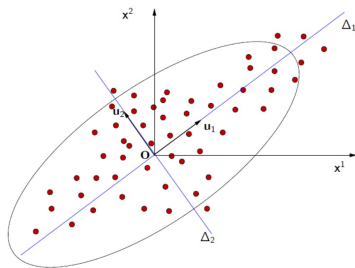★ Necessity to scale data (by their variances)

Data scaling

* We define metric $< ., . >_V$, with $V = diag(\sigma_1^2, \ldots, \sigma_N^2)$
* The scaled covariance matrix Σ_V is the **correlation matrix** R of initial data
* Spectral decomposition of R (instead of Σ)
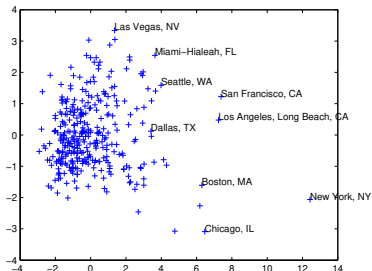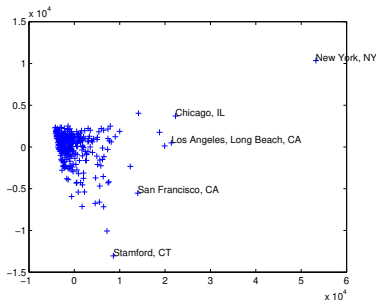
Optimality of PCA

Gaussian distribution

⋆ PCA decomposes the covariance matrix along its eigenvalues/vectors

⋆ Optimal basis to represent $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$

PCA is optimal to represent data whose distribution is close to Normal

Limitations

1. PCA considers the correlation between variables $\rightarrow$ linear relationships
   $\Rightarrow$ no non-linear relationship modeled
2. PCA optimises a quadratic loss
   $\Rightarrow$ sensitive to extreme values (outliers)
3. PCA is optmal for Gaussian distributed data
   $\Rightarrow$ not useful on clustered data